

Introduction to classification and regression trees, random forests and model-based recursive partitioning in R

Day 2: Stability and tree ensembles

Exercise 3: Glaucoma data revisited (2)

- a) Fit a random forest using the `cforest` function from package `party`, with either bootstrap sampling with size N or subsampling with size $.632 \cdot N$.

```
library(party)
data("GlaucomaM", package = "TH.data")
cfor_contr_bs <- cforest_control(replace = TRUE)
cfor_contr_ss <- cforest_control(replace = FALSE, fraction = 0.632)

set.seed(2908)
for_glauc_bs <- cforest(Class ~ ., data = GlaucomaM, control = cfor_contr_bs)
set.seed(21)
for_glauc_ss <- cforest(Class ~ ., data = GlaucomaM, control = cfor_contr_ss)
```

- b) Using the `predict()` function, compute the percentage of correctly classified observations — based on all trees and based only on those trees for which the observations are out-of-bag (by adding `OOB = FALSE` or `OOB = TRUE` to the function call) — and compare your results. Determine the predicted response class of each case and create a confusion matrix.

```
# calculate training errors:
#
# bootstrap sampling
predicted_class <- predict(for_glauc_bs)
table(GlaucomaM$Class, predicted_class)
```

```
##           predicted_class
##           glaucoma normal
## glaucoma           86     12
## normal              4     94
```

```
# subsampling
predicted_class <- predict(for_glauc_ss)
table(GlaucomaM$Class, predicted_class)
```

```
##           predicted_class
##           glaucoma normal
## glaucoma           84     14
## normal              5     93
```

```
sum(GlaucomaM$Class == predicted_class)/nrow(GlaucomaM)
```

```
## [1] 0.9030612
```

```
# OOB errors
#
# bootstrap sampling
predicted_class <- predict(for_glauc_bs, OOB = TRUE)
table(GlaucomaM$Class, predicted_class)
```

```
##           predicted_class
##           glaucoma normal
```

```
##      glaucoma      79      19
##      normal       15      83

#
# subsampling
predicted_class <- predict(for_glauc_ss, OOB = TRUE)
table(GlaucomaM$Class, predicted_class)
```

```
##           predicted_class
##           glaucoma normal
##      glaucoma      76      22
##      normal      10      88
```

```
sum(GlaucomaM$Class==predicted_class)/nrow(GlaucomaM)
```


```
## [1] 0.8367347
```

c) Compute and plot the standard permutation importance for the predictor variables.

```
set.seed(14)
imp <- varimp(for_glauc_bs)
imp[order(imp, decreasing = TRUE)]
```

```
##      vari      varg      vars      varn      tmg
## 2.473086e-02 2.133656e-02 1.256608e-02 1.136261e-02 8.918427e-03
##      phci      hic      tmi      mhci      rnf
## 8.809597e-03 7.904260e-03 7.788858e-03 7.385486e-03 7.355221e-03
##      phcg      tms      abri      mdic      abrs
## 6.945339e-03 6.389034e-03 6.312554e-03 5.436358e-03 5.119845e-03
##      vbri      mhcg      phcn      vart      hvc
## 4.888464e-03 4.000060e-03 3.968590e-03 3.657933e-03 3.077886e-03
##      emd      abrg      as      mhct      ag
## 2.682900e-03 2.475112e-03 1.977648e-03 1.475119e-03 1.395082e-03
##      vass      tmt      mhcs      vbrg      vbsi
## 1.287167e-03 1.218121e-03 1.204970e-03 1.172589e-03 1.122333e-03
##      mdn      tmn      ai      mdt      vasg
## 1.114135e-03 1.088222e-03 1.067874e-03 1.066057e-03 1.041898e-03
##      vbrs      at      eas      phcs      vbsg
## 1.039984e-03 1.021327e-03 9.937465e-04 8.209638e-04 8.047762e-04
##      abrn      phct      vbst      mr      an
## 7.697780e-04 7.635571e-04 6.894481e-04 6.429624e-04 6.317809e-04
##      vbrn      mhcn      vbss      mdi      vasi
## 5.379197e-04 5.239234e-04 5.154076e-04 5.045643e-04 5.044782e-04
##      eat      eag      abrt      eai      vast
## 4.305072e-04 3.883825e-04 3.855851e-04 3.524894e-04 3.418676e-04
##      vbrr      mdg      mds      vbsn      mv
## 3.186083e-04 1.548512e-04 1.379876e-04 6.698471e-05 -1.158468e-04
##      vasn      ean
## -1.915117e-04 -3.556480e-04
```

```
# plot importances for only a subset of the variables:
barplot(imp[42:49], names.arg = names(GlaucomaM[,42:49]), main = "variable importance")
```




```
set.seed(15)
# plot importances for only a subset of the variables:
imp <- varimp(for_glauc_ss)
imp[order(imp, decreasing = TRUE)]
```

```
##      vari      varg      vars      tmg      tmi
## 2.797222e-02 2.775000e-02 1.950000e-02 1.197222e-02 1.025000e-02
##      hic      tms      varn      phcg      phci
## 1.002778e-02 9.916667e-03 8.888889e-03 8.444444e-03 7.833333e-03
##      mhci      abri      rnf      mdic      vbri
## 6.500000e-03 6.361111e-03 6.333333e-03 5.166667e-03 4.333333e-03
##      phcn      abrg      abrs      vbsi      vbrs
## 3.638889e-03 3.277778e-03 3.250000e-03 3.083333e-03 2.583333e-03
##      hvc      mhcn      mhcg      emd      mhcs
## 2.277778e-03 2.138889e-03 2.111111e-03 1.972222e-03 1.861111e-03
##      vart      tmt      vasi      phcs      ai
## 1.861111e-03 1.833333e-03 1.444444e-03 1.361111e-03 1.250000e-03
##      mr      vasg      vass      an      vbrg
## 1.166667e-03 1.027778e-03 9.722222e-04 7.777778e-04 7.500000e-04
##      ean      as      mdn      tmn      abrn
## 6.944444e-04 6.666667e-04 6.388889e-04 6.388889e-04 6.111111e-04
##      vbsn      vbsg      abrt      ag      phct
## 5.833333e-04 5.555556e-04 4.722222e-04 4.722222e-04 4.722222e-04
##      eag      vbrn      mv      eai      vbst
## 4.722222e-04 3.888889e-04 3.055556e-04 2.777778e-04 2.500000e-04
##      vbss      mdt      at      mds      vast
```

```
## 2.500000e-04 2.500000e-04 2.500000e-04 1.944444e-04 1.944444e-04
##      eat      vbrt      eas      mdi      vasn
## 1.388889e-04 8.333333e-05 8.333333e-05 -5.551115e-20 -2.777778e-05
##      mdg      mhct
## -5.555556e-05 -5.555556e-05
```

```
barplot(imp[42:49], names.arg = names(GlaucomaM[,42:49]), main = "variable importance")
```



*# Note: Before interpreting any variable importances make sure that
your results are stable by means of setting two different
random seeds, before fitting the forest and computing the
variable importance. If the ranking of the top variable importances
changes between the two runs, this means that your forest was not
large enough (as compared to the number of variables) to generate
stable results and you should increase the number of trees in the
forest (ntree).*

d) Do the `stabletree()` function and `cforest()` function give the same 3 most important variables?
Set random seeds everywhere in the code where random sampling is involved.

*# According to the stablelearner analyses, vari, vasg and varg appear most often
in trees build on bootstrap datasets.
According to the random forest importances, vari, varg and vars are most
important.
The methods do not completely agree, but the number of times a variable appears
in trees grown on random samples of the data may not be the best indicator of
influence (importance), and can be expected to yield different results than*

```
# calculating variable importances in a random forest.
```

Exercise 4: Carseat sales revisited

Using the `cforest()` function from the `party` package, fit a random forest and a bagged ensemble to the 300 training observations. Using the `predict()` function, calculate the mean squared error (MSE) for:

$$MSE = \frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$$

a) The training observations

```
data("Carseats", package = "ISLR")
set.seed(42)
train <- sample(1:400, 300)
library(party)
bagged_cs <- cforest(Sales ~ ., data = Carseats[train,], controls = cforest_unbiased(mtry = NULL))
ranfor_cs <- cforest(Sales ~ ., data = Carseats[train,])

mean((Carseats[train, "Sales"] - predict(bagged_cs))^2)

## [1] 2.10367
mean((Carseats[train, "Sales"] - predict(ranfor_cs))^2)
```

```
## [1] 2.241822
```

b) The OOB predictions for training observations (add `OOB = TRUE` to the function call)

```
mean((Carseats[train, "Sales"] - predict(bagged_cs, OOB = TRUE))^2)

## [1] 3.204182
mean((Carseats[train, "Sales"] - predict(ranfor_cs, OOB = TRUE))^2)
```

```
## [1] 3.174919
```

c) The test observations (specify `newdata` argument)

```
mean((Carseats[-train, "Sales"] - predict(bagged_cs, newdata = Carseats[-train,]))^2)

## [1] 3.237652
mean((Carseats[-train, "Sales"] - predict(ranfor_cs, newdata = Carseats[-train,]))^2)
```

```
## [1] 3.389324
```

d) Are the OOB training error estimates indeed less optimistic / more realistic estimates of the test error?

```
# Yes, the error estimated based on OOB observations more closely resembles the error on test data
```

Exercise 5: Boston housing data

Fit a random forest and bagged ensemble of CART trees to the Boston housing data. This dataset reports the median value (`medv`) of owner-occupied homes in 506 districts in the Boston area (USA), together with 13 potential predictor variables reporting sociodemographic and other characteristics of the districts:

```
library("randomForest")
```

```
## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
?randomForest

## starting httpd help server ...

## done

data("Boston", package = "MASS")
?MASS::Boston
```

Use the `randomForest()` function from the `randomForest` package to predict the median value of owner occupied home. Fit both a random forest and a bagged ensemble. You could create the ensembles using a formula statement (i.e., `medv ~ .`), but if you supply arguments `x`, `y`, `xtest` and `ytest`, you can evaluate the test error for random forests consisting of 1 through the specified number of trees to grow. You will use these estimates to evaluate the effect of the number of trees on accuracy and to compare the OOB error estimates with the error estimated on test data. Therefore, first separate the data into a training and test set:


```
set.seed(1)
train <- sample(1:nrow(Boston), nrow(Boston)/4)
x <- Boston[train, -14]
y <- Boston[train, 14]
xtest <- Boston[-train, -14]
ytest <- Boston[-train, 14]
```

- a) Set the `ntree` argument to 750, and create two ensembles: a bagged ensemble, using all 13 predictor variables for selecting each split, and a random forest using `sqrt(13)` for selecting each split.

```
set.seed(1)
rf.boston <- randomForest(x = x, y = y, xtest = xtest, ytest = ytest, mtry = sqrt(13),
                          ntree = 750, data = Boston)
bag.boston <- randomForest(x = x, y = y, xtest = xtest, ytest = ytest, mtry = 13,
                          ntree = 750)
```

- b) Inspect the test error (which can be found in the `$test$mse` slot of the `randomForest` object) and the OOB error (which can be found in the `$mse` slot of the `randomForest` object) for both ensembles. Create a plot where the number of trees in the ensemble is on the x-axis, and the MSE is on the y-axis. Draw different lines for the OOB and test error of the bagged ensemble, as well as the OOB and test error of the random forest. (Hint: use functions `plot()` and `lines()` and the `type`, `lty` and `col` arguments)

```
plot(1:750, rf.boston$test$mse, col = "red", type = "l", xlab = "Number of Trees",
     ylab = "Test MSE")
lines(1:750, rf.boston$mse, col = "red", type = "l", lty = 2)
lines(1:750, bag.boston$test$mse, col = "blue", type = "l")
lines(1:750, bag.boston$mse, col = "blue", type = "l", lty = 2)
legend("topright", paste(rep(c("rf", "bagged"), times=2),
                        rep(c("(test)", "(OOB)"), each = 2)),
      col = rep(c("red", "blue"), times = 2), cex = .5, lty = rep(1:2, each = 2))
```



c) After which number of trees does the MSE start to stabilize?

After about 200 trees.

d) Which of the two ensembles is most accurate on the test data?

If we would use 600-750 trees for prediction, the bagged and random forest ensemble seem equally accurate.

e) Does the OOB error give a better estimate of test error for the random forest, or for the bagged ensemble? Can you explain the difference?

The OOB error gives an estimate of prediction error that is closer to the test error for the bagged ensemble (at least, in this case it does).

#

The OOB error overestimates the test error more for the random forest than for the bagged ensemble. This may be because the individual trees in the random forest may be somewhat less accurate, because they cannot always pick up the best predictor variables (due to the random selection of predictor variables as splitting candidates). As a result, the OOB error may be somewhat higher (less optimistic) for random forests than for bagged ensembles.